

# Graph Random Neural Networks for Semi-Supervised Learning on Graphs

Presented by:

Gauransh Sawhney 2018A3PS0325P

Utkarsh Kumar Singh 2018A3PS0368P

In partial fulfillment of the requirements of the course:


BITS F464 Machine Learning

Submitted to: Dr. Kamlesh Tiwari


# Semi-Supervised Learning

- What is Semi-Supervised Learning
- Semi-Supervised Learning on Graphs

# Semi-Supervised Learning on Graphs



**Input:** a partially labeled & attributed graph





**Output:** infer the labels of unlabeled nodes

# Graph Neural Networks

- What are Graph Neural Networks(GNNs)
- Problems with existing implementations

# Graph Neural Network



node  $v$ 's embedding at  $k + 1$

non-linear activation function (e.g. ReLU)

$$H^{k+1} = \sigma(\hat{A}H^{(k)}W^{(k)})$$

normalized Laplacian matrix

$$H^{k+1} = \sigma \left( W^k \sum_{u \in N(v) \cup v} \frac{H_u^k}{\sqrt{|N(u)||N(v)|}} \right)$$

the neighbors of node  $v$

# Existing Issues

$$H^{k+1} = \sigma(\hat{A}H^{(k)}W^{(k)})$$

- Oversmoothing
  - Stacking multiple GNN layers makes nodes indistinguishable; coupling the **feature propagation** and **non-linear transformation** steps, aggravates this problem
- Not robust to graph attacks
  - Each node is highly dependent on neighbors, making it non-robust to noise
- Overfitting in case of semi-supervised
  - In the standard setting of semi-supervised training, scarce node label information can be overfit



# GRAND: Graph Random Neural Network

- Architecture
- Algorithm
- How does it tackle the issues faced by other GNNs



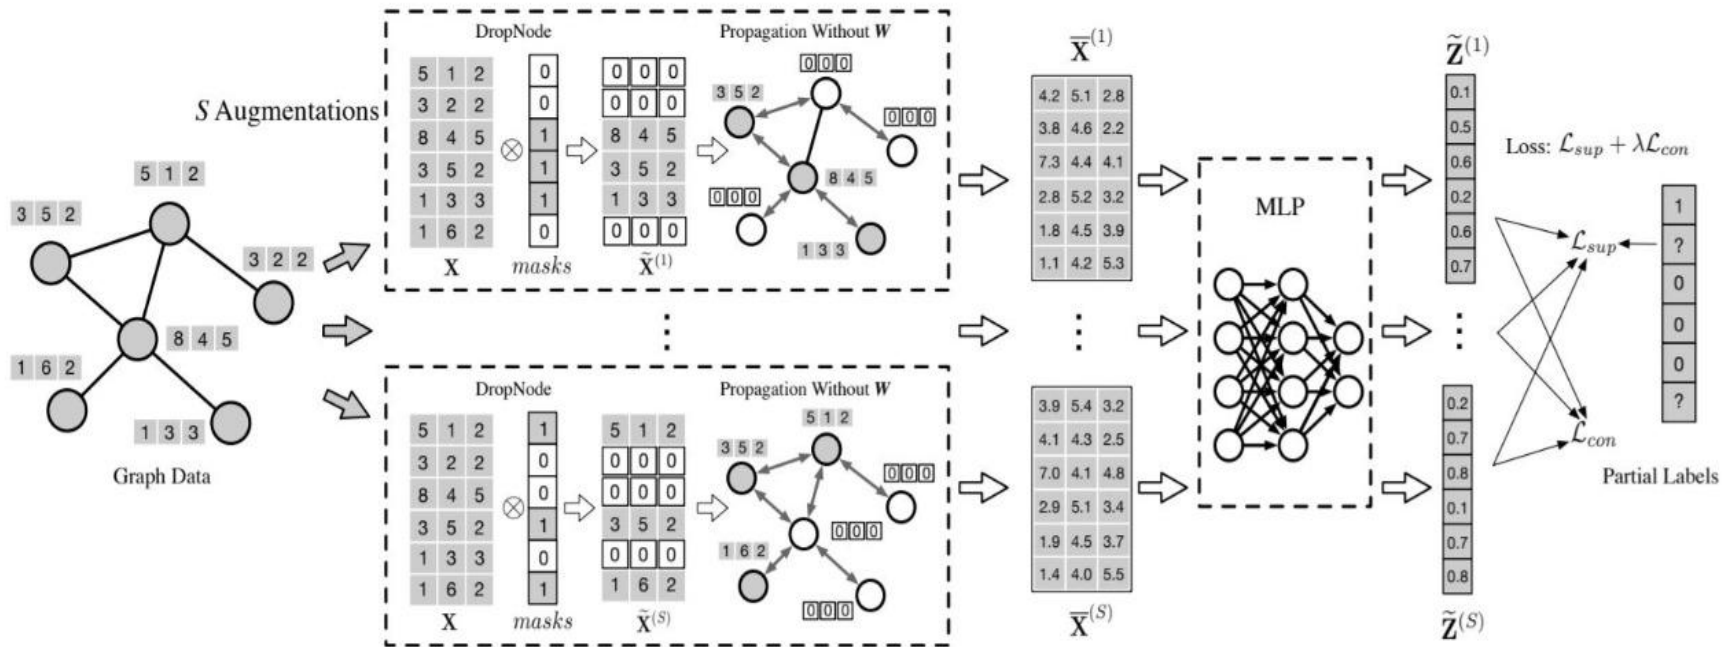
# GRAND

## Random Propagation





# GRAND





# Algorithm

---

## Algorithm 1 GRAND

---

### Input:

Adjacency matrix  $\hat{\mathbf{A}}$ , feature matrix  $\mathbf{X} \in \mathbb{R}^{n \times d}$ , times of augmentations in each epoch  $S$ , DropNode/dropout probability  $\delta$ , learning rate  $\eta$ , an MLP model:  $f_{mlp}(\mathbf{X}, \Theta)$ .

### Output:

Prediction  $\mathbf{Z}$ .

- 1: **while** not convergence **do**
  - 2:   **for**  $s = 1 : S$  **do**
  - 3:     Pertube the input:  $\tilde{\mathbf{X}}^{(s)} \sim \text{DropNode}(\mathbf{X}, \delta)$ .
  - 4:     Perform propagation:  $\overline{\mathbf{X}}^{(s)} = \frac{1}{K+1} \sum_{k=0}^K \hat{\mathbf{A}}^k \tilde{\mathbf{X}}^{(s)}$ .
  - 5:     Predict class distribution using MLP:  $\tilde{\mathbf{Z}}^{(s)} = f_{mlp}(\overline{\mathbf{X}}^{(s)}, \Theta)$
  - 6:   **end for**
  - 7:   Compute supervised classification loss  $\mathcal{L}_{sup}$  via Eq. 1 and consistency regularization loss via Eq. 3.
  - 8:   Update the parameters  $\Theta$  by gradients descending:  $\Theta = \Theta - \eta \nabla_{\Theta} (\mathcal{L}_{sup} + \lambda \mathcal{L}_{con})$
  - 9: **end while**
  - 10: Output prediction  $\mathbf{Z}$  via:  $\mathbf{Z} = f_{mlp}(\frac{1}{K+1} \sum_{k=0}^K \hat{\mathbf{A}}^k \mathbf{X}, \Theta)$ .
-



# Loss Functions

$$\mathcal{L}_{sup} = -\frac{1}{S} \sum_{s=1}^S \sum_{i=0}^{m-1} \mathbf{Y}_i^\top \log \tilde{\mathbf{Z}}_i^{(s)}.$$

$$\bar{\mathbf{z}}'_{ij} = \bar{\mathbf{z}}_{ij}^{\frac{1}{T}} \bigg/ \sum_{c=0}^{C-1} \bar{\mathbf{z}}_{ic}^{\frac{1}{T}}, (0 \leq j \leq C-1),$$

$$\mathcal{L}_{con} = \frac{1}{S} \sum_{s=1}^S \sum_{i=0}^{n-1} \|\bar{\mathbf{z}}'_i - \tilde{\mathbf{Z}}_i^{(s)}\|_2^2.$$

# Results

Some of the results presented in the paper:

- Comparison with existing architectures on benchmarks
- Generalization analysis
- Robustness analysis
- Over-smoothing analysis
- Results on large datasets



# Dataset Description

**3 datasets were used to benchmark results**

Dataset	Nodes	Edges	Train/Valid/Test Nodes	Classes	Features	Default Label Rate
Cora	2708	5429	140/500/1000	7	1433	0.052
Citeseer	3327	4732	120/500/1000	6	3703	0.036
Pubmed	19717	44338	60/500/1000	3	500	0.003


# Comparison with existing architectures

Method	Cora	Citeseer	Pubmed
GCN [20]	81.5	70.3	79.0
GAT [35]	83.0±0.7	72.5±0.7	79.0±0.3
APPNP [21]	83.8±0.3	71.6± 0.5	79.7 ± 0.3
Graph U-Net [12]	84.4±0.6	73.2±0.5	79.6±0.2
SGC [39]	81.0 ±0.0	71.9 ± 0.1	78.9 ± 0.0
MixHop [1]	81.9± 0.4	71.4±0.8	80.8±0.6
GMNN [31]	83.7	72.9	81.8
GraphNAS [13]	84.2±1.0	73.1±0.9	79.6±0.4
GraphSAGE [17]	78.9±0.8	67.4±0.7	77.8±0.6
FastGCN [7]	81.4±0.5	68.8±0.9	77.6±0.5
VBAT [9]	83.6±0.5	74.0±0.6	79.9±0.4
G <sup>3</sup> NN [25]	82.5±0.2	74.4±0.3	77.9 ±0.4
GraphMix [36]	83.9±0.6	74.5±0.6	81.0±0.6
DropEdge [32]	82.8	72.3	79.6
GRAND_dropout	84.9±0.4	75.0±0.3	81.7±1.0
GRAND_DropEdge	84.5±0.3	74.4±0.4	80.9±0.9
GRAND_GCN	84.5±0.3	74.2±0.3	80.0±0.3
GRAND_GAT	84.3±0.4	73.2± 0.4	79.2±0.6
GRAND	<b>85.4±0.4</b>	<b>75.4±0.4</b>	<b>82.7±0.6</b>
w/o CR	84.4±0.5	73.1±0.6	80.9±0.8
w/o mDN	84.7±0.4	74.8±0.4	81.0±1.1
w/o sharpening	84.6±0.4	72.2±0.6	81.6±0.8
w/o CR & DN	83.2±0.5	70.3±0.6	78.5±1.4


Table 1: Overall classification accuracy (%).




## Generalization Analysis



(a) Without RP




(b) Without CR

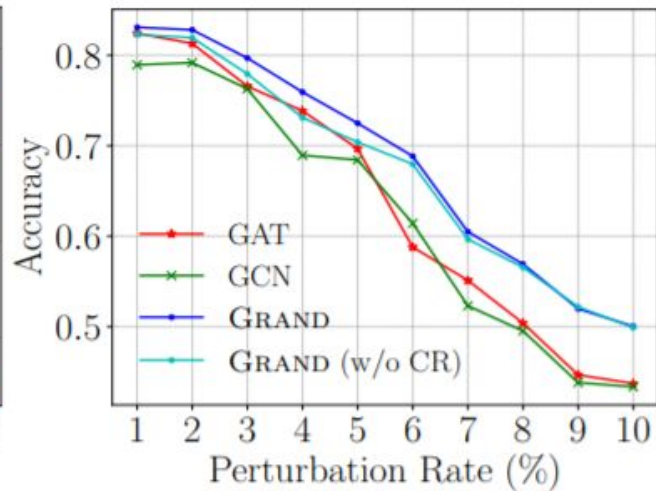


(c) GRAND(with RP and CR)

# Robustness Analysis




(a) Random Attack



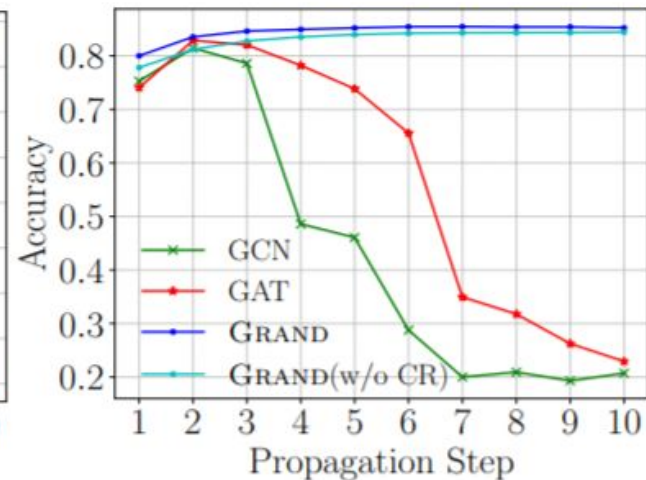
(b) Metattack



# Over-smoothing analysis



(a) MADGap

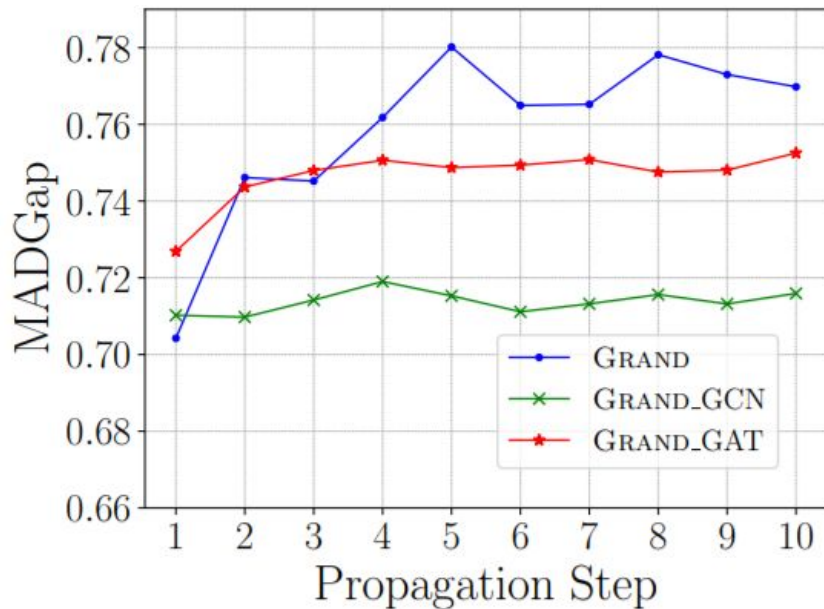


(b) Classification Result



## Other results presented in the paper

### Over-smoothness of GRAND and its variants(on Cora)





## Other results presented in the paper

### Classification Accuracy of GRAND on large datasets

Method	Cora Full	Coauthor CS	Coauthor Physics	Amazon Computer	Amazon Photo	Aminer CS
GCN	$62.2 \pm 0.6$	$91.1 \pm 0.5$	$92.8 \pm 1.0$	$82.6 \pm 2.4$	$91.2 \pm 1.2$	$49.9 \pm 2.0$
GAT	$51.9 \pm 1.5$	$90.5 \pm 0.6$	$92.5 \pm 0.9$	$78.0 \pm 19.0$	$85.7 \pm 20.3$	$49.6 \pm 1.7$
GRAND	<b><math>63.5 \pm 0.6</math></b>	<b><math>92.9 \pm 0.5</math></b>	<b><math>94.6 \pm 0.5</math></b>	<b><math>85.7 \pm 1.8</math></b>	<b><math>92.5 \pm 1.7</math></b>	<b><math>52.8 \pm 1.2</math></b>

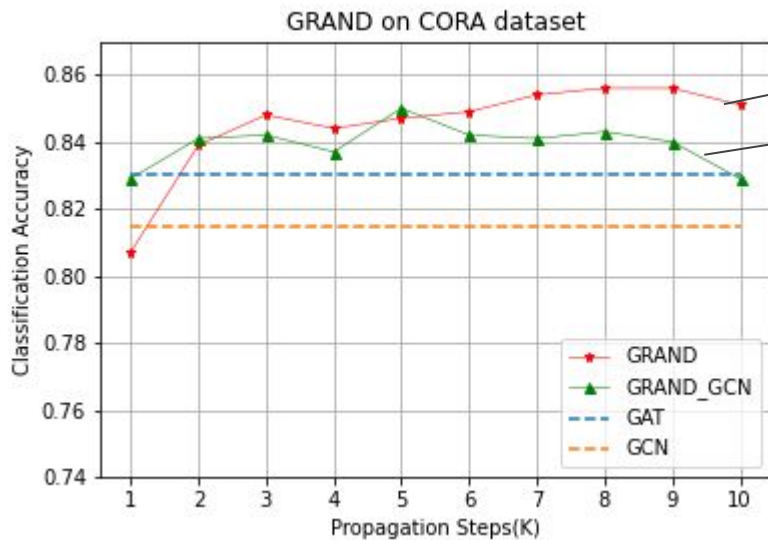
# Experiments

The following experiments were conducted:

- MLP v/s GCN as classification network
- Classification accuracy v/s  $\{K, S\}$
- Sensitivity wrt CR loss coefficient  $\lambda$

# 1. MLP v/s GCN


## Effect of using an MLP vs GCN as the classification network



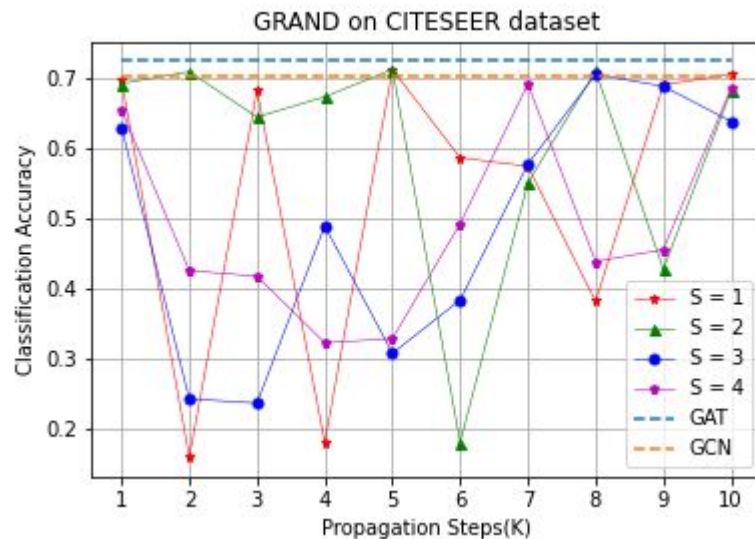
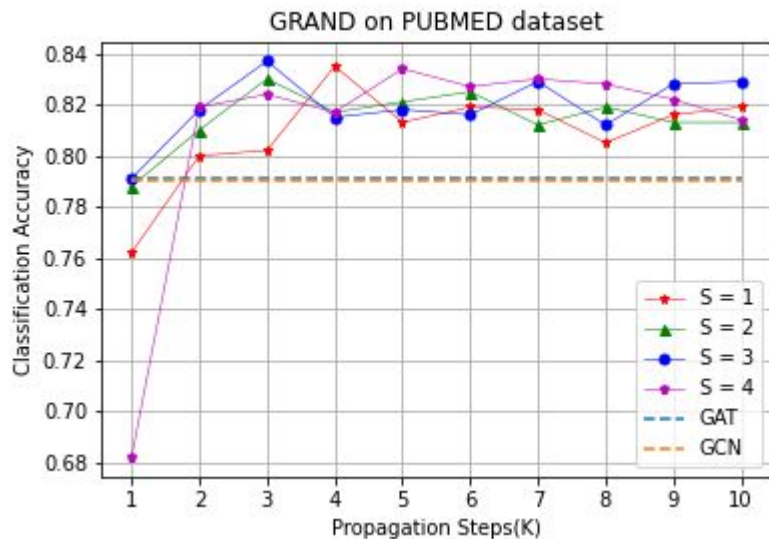
**GRAND** very clearly outperforms **GRAND\_GCn** in terms of *classification accuracy*

## 2. Classification Accuracy v/s {K, S}

(i) Effect of K(propagation order) and S(number of data augmentations) on Classification Accuracy on **GRAND**(DropNode data augmentation)





## 2. Classification Accuracy v/s {K, S}



## 2. Classification Accuracy v/s {K, S}

(ii) Effect of K(propagation order) and S(number of data augmentations) on Classification Accuracy on **GRAND\_dropout** and **GRAND\_dropedge**(alternative data augmentation techniques)






### 3. Sensitivity wrt $\lambda$

#### Classification Accuracy v/s $\lambda$

Consistency Regularization  
Loss Coefficient




MLP classification network

Both using DropNode for  
data augmentation

GCN classification network

### 3. Sensitivity wrt $\lambda$

#### Classification Accuracy v/s $\lambda$

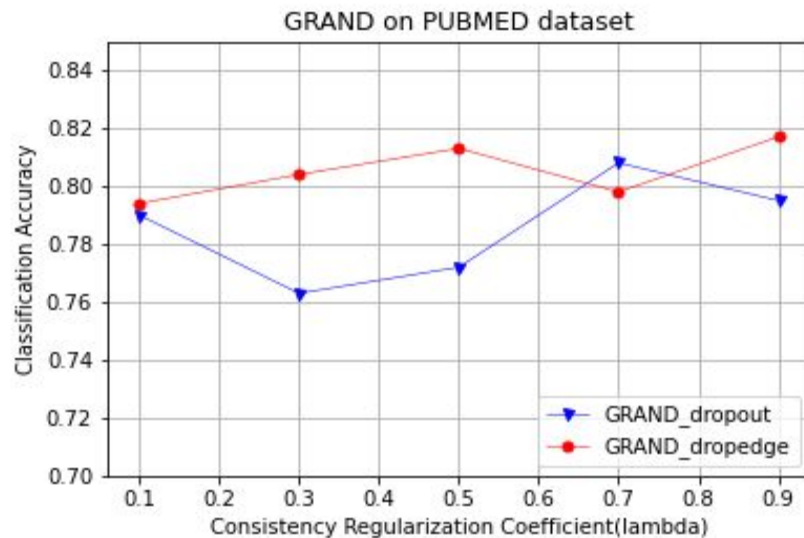
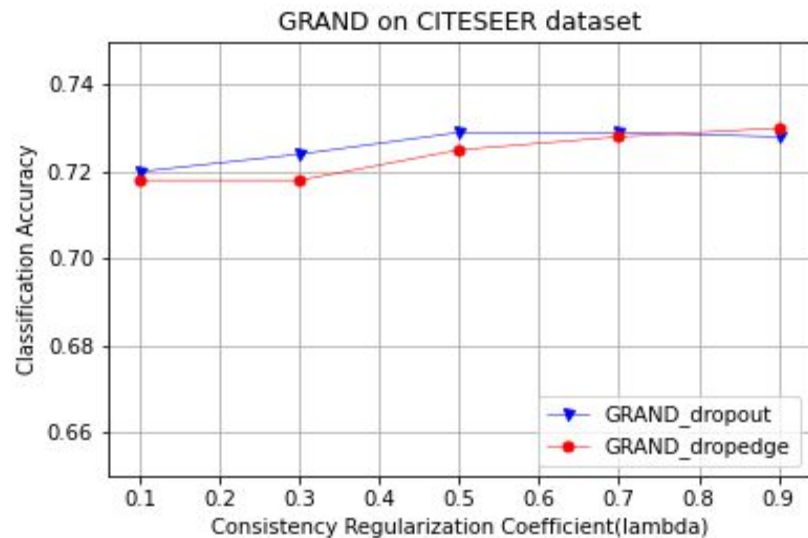


DropEdge data augmentation

Both using MLPs as  
classification networks

Dropout data augmentation

### 3. Sensitivity wrt $\lambda$



# Ablation Study

The effect of the absence of the following parameters was studied:

- w/o consistency regularization(CR)
- w/o multiple dropnode(mDN)
- w/o sharpening
- w/o consistency regularization(CR) and dropnode(DN)



# Ablation Study

these are classification accuracies

Method	Cora	Pubmed	Citeseer
w/o CR ( $\lambda=0$ )	0.841	0.811	0.728
w/o mDN ( $S=1$ )	0.85	0.80	0.744
w/o sharpening ( $T=1$ )	0.844	0.816	0.578
w/o CR & DN ( $\lambda=0, \delta=0$ )	0.835	0.787	0.597

**END**

**Thank You!**